

Discrete Optimization

A random-key genetic algorithm for the generalized traveling salesman problem

Lawrence V. Snyder ^{a,*}, Mark S. Daskin ^b

^a *Department of Industrial and Systems Engineering, Lehigh University, 200 West Packer Avenue, Mohler Lab, Bethlehem, PA 18015, USA*

^b *Department of Industrial Engineering and Management Sciences, Northwestern University, 2145 Sheridan Road, MEAS C210, Evanston, IL 60208, USA*

Received 18 April 2001; accepted 6 September 2004

Available online 17 May 2005

Abstract

The generalized traveling salesman problem is a variation of the well-known traveling salesman problem in which the set of nodes is divided into clusters; the objective is to find a minimum-cost tour passing through one node from each cluster. We present an effective heuristic for this problem. The method combines a genetic algorithm (GA) with a local tour improvement heuristic. Solutions are encoded using random keys, which circumvent the feasibility problems encountered when using traditional GA encodings. On a set of 41 standard test problems with symmetric distances and up to 442 nodes, the heuristic found solutions that were optimal in most cases and were within 1% of optimality in all but the largest problems, with computation times generally within 10 seconds. The heuristic is competitive with other heuristics published to date in both solution quality and computation time.

© 2005 Elsevier B.V. All rights reserved.

Keywords: Traveling salesman; GTSP; Genetic algorithms; Random keys; Metaheuristics

1. Introduction

1.1. The GTSP

The *generalized traveling salesman problem* (GTSP) is a variation of the well-known traveling salesman problem in which not all nodes need to be visited by the tour. In particular, the set V of nodes is partitioned into m sets, or clusters,

* Corresponding author. Tel.: +1 610 758 6696; fax: +1 610 758 4886.

E-mail addresses: larry.snyder@lehigh.edu (L.V. Snyder), m-daskin@northwestern.edu (M.S. Daskin).

V_1, \dots, V_m with $V_1 \cup \dots \cup V_m = V$ and $V_i \cap V_j = \emptyset$ if $i \neq j$. The objective is to find a minimum-length tour containing exactly one node from each set V_i .

Applications of the GTSP include postal routing [19], computer file sequencing [16], order picking in warehouses [27], process planning for rotational parts [5], and the routing of clients through welfare agencies [35]. In addition, many other combinatorial optimization problems can be reduced to, or have as a subproblem, the GTSP (see [19]). The GTSP is NP-hard, since the TSP is a special case obtained by partitioning V into $|V|$ subsets, each containing one node. In this paper we present a heuristic approach to solving the GTSP.

A given instance of the GTSP may use “symmetric” distances, in which the distance from node i to node j is the same as the distance from j to i , or “asymmetric” distances; there are theoretical and computational consequences of using one type of distance metric or the other. For example, the GTSP is often transformed into an equivalent TSP, but many such transformations depend on whether or not the problem is symmetric. Similarly, a given algorithm’s computational performance may be quite different when applied to different types of instances. We test our heuristic on a standard test-bed of symmetric instances.

1.2. Genetic algorithms

A genetic algorithm (GA) is a metaheuristic inspired by the efficiency of natural selection in biological evolution. Genetic algorithms have been applied successfully to a wide variety of combinatorial optimization problems and are the subject of numerous recent books [26,30,39–41] and conference proceedings [1,2,4,17]. Unlike traditional heuristics (and some metaheuristics like tabu search) that generate a single solution and work hard to improve it, GAs maintain a large number of solutions and perform comparatively little work on each one. The collection of solutions currently under consideration is called the *population*. Each member of the population (called an *individual* or a *chromosome*) is an encoded version of a solution. The encoding strategy is different for different

optimization problems, and a given problem may have more than one workable encoding strategy. The goal of any encoding is to translate a solution into a string of *genes* that make up the chromosome, just as in biological genetics.

Each iteration of a GA consists of several *operators* that construct a new *generation* of solutions from the old one in a manner designed to preserve the genetic material of the better solutions (survival of the fittest). Many GA operators have been proposed; the three most common are reproduction, crossover, and mutation. *Reproduction* consists of simply copying the best solutions from the previous generation into the next, with the intention of preserving very high-quality solutions in the population as-is. *Crossover* takes two “parents,” randomly chosen, and produces one or more “offspring” that contain some combination of genes from the parents. Crossover can be performed in a deterministic manner (e.g., “one-point” crossover), with genes appearing before a certain cutoff coming from parent 1 and genes after the cutoff coming from parent 2, or in a random manner, with each gene taken from a given parent with a certain probability. The *mutation* operator changes a few genes randomly, borrowing from the evolutionary concept that random genetic mutations may produce superior offspring (or, of course, inferior offspring, but such individuals are less likely to survive from one generation to the next). Our algorithm does not use mutation, but rather *immigration*, in which new individuals are generated randomly from scratch, rather than performing random mutations on existing individuals.

Each iteration of a standard GA consists of generating a new population using the GA operators. Our GA also uses judiciously applied local improvement heuristics to improve the quality of the individuals in the population without adding excessive computation time.

1.3. Random keys

The GA presented in this paper uses *random keys* to encode solutions. The use of random keys is described in [3] and is useful for problems that require permutations of the integers and for which

traditional one- or two-point crossover presents feasibility problems. The technique is best illustrated with an example.

Consider a 5-node instance of the classical TSP. Traditional GA encodings of TSP solutions consist of a stream of integers representing the order in which nodes are to be visited by the tour. (That is, the solution 4 2 1 5 3 represents the tour $3 \rightarrow 2 \rightarrow 5 \rightarrow 1 \rightarrow 4$, not the tour $4 \rightarrow 2 \rightarrow 1 \rightarrow 5 \rightarrow 3$.) But one-point crossover, for example, may result in multiple nodes assigned to the same slot on the tour. For example, the parents

1	4	3	2	5
2	4	5	3	1

might produce the children

1	4	3	3	1
2	4	5	2	5

In the random key method, we assign each gene a random number drawn uniformly from $[0, 1)$. To decode the chromosome, we visit the nodes in ascending order of their genes. For example,

Random key	0.42	0.06	0.38	0.48	0.81
Decodes as	3	1	2	4	5

(Ties are broken arbitrarily.) Nodes that should be early in the tour tend to “evolve” genes closer to 0 and those that should come later tend to evolve genes closer to 1. Standard crossover techniques will generate children that are guaranteed to be feasible.

1.4. Contributions

This paper contributes to the literature by presenting a heuristic for the GTSP that uses a novel combination of features (GAs, random keys, and improvement heuristics). Our heuristic finds optimal solutions for a large percentage of instances from a set of standard test problems, with computation times that are competitive with, if not superior to, previously published heuristics. We also present extensive computational results lending insight into the behavior of the heuristic, such as the relative importance of the GA and the improvement heuristics. Finally, like any GA,

our heuristic is simple to implement and can be easily modified to incorporate alternate objective functions or additional constraints.

The remainder of this paper is organized as follows: Section 2 presents a brief literature review; Section 3 describes the algorithm in greater detail; Section 4 provides computational results; and Section 5 summarizes our findings and discusses avenues for future research.

2. Literature review

The earliest papers on the GTSP discuss the problem in the context of particular applications [16,35,38]. Laporte et al. [20,21] and Laporte and Nobert [22] began to study exact algorithms and some theoretical aspects of the problem in the 1980s. Since then, a few dozen papers related to the GTSP have appeared in the literature. Fischetti et al. [13] discuss facet-defining inequalities for the GTSP polytope, and in a later paper [14] use the polyhedral results to develop a branch-and-cut algorithm. Other exact algorithms are presented in [28] (a Lagrangian-based branch-and-bound algorithm) and [7] (a dynamic programming algorithm).

A variety of descriptions of the GTSP are present in the literature. Some papers assume symmetry or the triangle inequality; others do not. Some require the subsets to form a strict partition of the node set; others allow them to overlap. Most papers allow the tour to visit more than one node per cluster, but some require that exactly one node per cluster be visited. (The two formulations are equivalent if the distance matrix satisfies the triangle inequality.) Only a few papers (see [22,35]) handle fixed costs for including a node on the tour, possibly because such fixed costs can be incorporated into the distance costs via a simple transformation (see [22]; note that this transformation does not preserve the triangle inequality).

Applications of the GTSP are discussed in [5,19,27]. Ref. [12] transforms the clustered rural postman problem into the GTSP. Refs. [6,9,23,25,29] present transformations of the GTSP into standard TSP instances that can then be solved using TSP algorithms and heuristics. Some of the

resulting TSP instances have nearly the same number of nodes as the original GTSP instances and have nice properties like symmetry and the triangle inequality; others have many more nodes and are more irregular. Some transformations of the GTSP into the TSP (e.g., that in [29]) have the property that an optimal solution to the related TSP can be converted to an optimal solution to the GTSP, but a (sub-optimal) feasible solution for the TSP may not be feasible for the GTSP, reinforcing the necessity for good heuristics for the GTSP itself. Furthermore, well-known heuristics for the TSP may not perform well for the GTSP; for example, “generalized” versions of Christofides’s heuristic have worst-case bounds that are strictly greater than its worst-case bound of $\frac{3}{2}$ for the TSP [11,36].

Two approximation algorithms have been published for the GTSP. Slavík [36] develops a $3\rho/2$ -approximation algorithm for the GTSP, where ρ is the number of nodes in the largest cluster, that is, $\rho = \max_{i=1,\dots,m} \{|V_i|\}$. As ρ may be quite large, the worst-case bound may be relatively weak. Garg et al. [15] present an approximation algorithm for the group Steiner tree problem, which provides as a byproduct an $O(\log^2 |V| \log \log |V| \log m)$ -approximation algorithm for the GTSP. Both papers assume that the distance matrix satisfies the triangle inequality.

Noon [27] proposed several heuristics for the GTSP, the most promising of which is an adaptation of the well-known nearest-neighbor heuristic for the TSP. Fischetti et al. [14] propose similar adaptations of the farthest-insertion, nearest-insertion, and cheapest-insertion heuristics.

Renaud and Boctor [32] develop the most sophisticated heuristic published to date, which they call GI^3 (Generalized Initialization, Insertion, and Improvement). GI^3 is a generalization of the I^3 heuristic proposed in [33]. It consists of three phases. The first, Initialization, constructs an initial tour by choosing one node from each cluster that is “close” to the other clusters, then greedily building a tour that passes through some, but not necessarily all, of the chosen nodes. The second phase, Insertion, completes the tour by successively inserting nodes from unvisited clusters in the cheapest possible manner between two consecutive

clusters on the tour, allowing the visited node to change for the adjacent clusters; after each insertion, the heuristic performs a modification of the 3-opt improvement method. The third phase, Improvement, uses modifications of 2-opt and 3-opt to improve the tour. The modifications, called G2-opt, G3-opt, and G-opt, allow the visited nodes from each cluster to change as the tour is being re-ordered by the 2-opt or 3-opt procedures.

Several researchers (see [18] and the references contained within) have implemented GAs for the standard TSP, with mixed results. The GA in [18] found new best solutions for some well-studied benchmark problems.

3. The heuristic

Our heuristic couples a GA much like that described in [3] with a local improvement heuristic. We will describe the heuristic in the context of the symmetric GTSP; the heuristic can be modified easily to handle problems with asymmetric distance matrices, though our computational results do not necessarily generalize to the asymmetric case. Throughout the rest of this paper, we will use the terms “solution,” “individual,” and “chromosome” interchangeably to refer to either a GTSP tour or its representation in the GA.

3.1. Encoding and decoding

Bean [3] suggests encoding the GTSP as follows: each set V_i has a gene consisting of an integer part (drawn from $\{1, \dots, |V_i|\}$) and a fractional part (drawn from $[0, 1)$). The integer part indicates which node from the cluster is included on the tour, and the nodes are sorted by their fractional part as described in Section 1.3 to indicate the order.

For example, consider a GTSP instance with $V = \{1, \dots, 20\}$ and $V_1 = \{1, \dots, 5\}, \dots, V_4 = \{16, \dots, 20\}$. The random key encoding

4.23 3.07 1.80 3.76

decodes as the tour $8 \rightarrow 4 \rightarrow 18 \rightarrow 11$: the integer parts of the genes decode as, respectively, node 4, node 8 (the third node in V_2), node 11 (the first node in V_3), and node 18 (the third node in V_4),

and the fractional parts, when sorted, indicate that the clusters should be visited in the order $2 \rightarrow 1 \rightarrow 4 \rightarrow 3$.

The GA operators act directly on the random key encoding. When tour improvements are made, the encoding must be adjusted to account for the new solution (see Section 3.4).

The population is always stored in order of objective value, from best to worst.

3.2. Initial population

The initial population is created by generating a population consisting of N chromosomes (we use $N = 100$), drawing the gene for set V_i uniformly from $[1, |V_i| + 1)$ (or, equivalently, adding an integer drawn randomly from $\{1, \dots, |V_i|\}$ to a real number drawn randomly from $[0, 1)$). “Level-I” improvement heuristics (see Section 3.4) are applied to each individual in the population.

3.3. GA operators

At each generation, 20% of the population comes directly from the previous population via reproduction; 70% is spawned via crossover; and the remaining 10% is generated via immigration. We describe each of these operators next.

3.3.1. Reproduction

Our algorithm uses an *elitist strategy* of copying the best solutions in the population to the next generation. This guarantees monotonic non-degradation of the best solution from one generation to the next and ensures a constant supply of good individuals for mating.

3.3.2. Crossover

We use *parametrized uniform crossover* [37] to generate offspring. First, two parents are chosen at random from the old population. One child is generated from the two parents, and it inherits each gene from parent 1 with probability 0.7 and from parent 2 with probability 0.3.

3.3.3. Immigration

A small number of new individuals are created in each generation using the technique described

in Section 3.2. Like the more typical mutation operator, this immigration operator helps ensure a diverse population.

3.4. Improvement heuristics

Local improvement heuristics have been shown by various researchers to add a great deal of power to GAs (see, for example, [24]). In our algorithm, every time a new individual is created, either during crossover or immigration or during the construction of the initial population, we attempt to improve the solution using 2-opt and “swap” operations.

The well-known 2-opt heuristic attempts to find two edges of the tour that can be removed and two edges that can be inserted to obtain a single tour with lower cost. For problems with Euclidean distances, a 2-opt exchange “straightens” a tour that crosses itself.

The “swap” operation involves removing the node from cluster V_i and inserting a different node from V_i into the tour. The insertion is done using a modified nearest-neighbor criterion, so that the new node may be inserted on the tour in a spot different from the original node. In pseudocode, the swap operation is as follows:

```
procedure SWAP(tour  $T$ ; set  $V_i$ ; node  $j \in V_i$ ,  

 $j \in T$ ; distances  $D_{uv}$  between each  $u, v \in V$ )  

  remove  $j$  from  $T$   

  for each  $k \in V_i$   

     $c_k \leftarrow \min\{D_{uk} + D_{kv} - D_{uv} \mid (u, v) \text{ is an edge in } T\}$   

 $k^* \leftarrow \operatorname{argmin}_{k \in V_i} \{c_k\}$   

  insert  $k^*$  into  $T$  between the nodes  $u$  and  $v$   

  that attained the cost  $c_{k^*}$ 
```

Note that k^* may be the same node as j , which allows the heuristic to move a node to a different spot on the tour.

The 2-opt and swap operations are considered separately by our heuristic; they are not integrated, as in the G2-opt, G3-opt, and G-opt methods used by [32].

When an improvement is made to the solution, the encoding must be altered to reflect the new tour. We do this by rearranging the fractional

parts of the genes to give rise to the new tour. When a swap is performed, we must also adjust the integer part of the gene for the affected set.

To take a simple example, suppose we have the chromosome

1.1 1.3 1.4 1.7

for the sets V_i given in Section 3.1; this represents the tour $1 \rightarrow 6 \rightarrow 11 \rightarrow 16$. Now suppose we perform a swap by replacing node 16 with node 17 and repositioning it after node 1 so that the new tour is $1 \rightarrow 17 \rightarrow 6 \rightarrow 11$. The new encoding is

1.1 1.4 1.7 2.3

For each individual in the population, we store both the original (pre-improvement) cost and the final cost after improvements have been made. When a new individual is created, we compare its pre-improvement cost to the pre-improvement cost of the individual at position $p \times N$ in the previous (sorted) population, where $p \in [0, 1]$ is a parameter of the algorithm (we use $p = 0.05$ in our implementation). If the new solution is worse than the pre-improvement cost of this individual we use *level-I improvement*: apply one 2-opt exchange and one swap (assuming a profitable one can be found) and store the resulting individual. If, on the other hand, the new solution is better, we use *level-II improvement*: apply 2-opts until no profitable 2-opts can be found, then apply swaps until no profitable swaps can be found, and repeat until no improvements have been made in a given pass. This technique is designed to spend more time improving solutions that seem promising in comparison to previous solutions and to spend less time on the others. In both level-I and level-II improvement, we use a “first-improving” strategy in which the first move of a given type that improves the objective value is implemented, rather than searching for the best such move before choosing one.

No improvement is done during the reproduction phase, since these solutions have already been improved.

3.5. Population management

Two issues are worth mentioning regarding how we maintain GA populations. First, no dupli-

cates are allowed in the population. Each time a new individual is created, it is compared to the individuals already in the population; if it matches one of them, it is discarded and a new individual is created. Duplicate-checking is performed *after* improvement, and individuals are considered identical if the *decoded solutions* are identical, not the genes. That is, chromosomes 2.14 4.25 3.50 2.68 and 2.07 4.73 3.81 2.99 are considered identical since they represent the same tour.

Second, it is undesirable to allow multiple “reflections” and “rotations” of a tour to co-exist in the population. That is, if $1 \rightarrow 6 \rightarrow 11 \rightarrow 16$ is in the population, we would not want $6 \rightarrow 11 \rightarrow 16 \rightarrow 1$ or $16 \rightarrow 11 \rightarrow 6 \rightarrow 1$ in the population. There are two reasons for this. One is that such quasi-duplicates appear to add diversity to the population but in fact do not, and so should be avoided just as duplicates are. The second reason is that crossover between two such individuals will result in offspring whose genetic information is “watered down” and may lead to slower convergence of the GA. In the course of the GA, nodes that should appear consecutively in the tour tend to evolve fractional parts of the gene that are close to one another. For this process to work, the beginning of the tour needs to be well defined—that is, we *can’t* allow reflections and rotations.

For example, suppose $V = \{1, \dots, 5\}$ and $V_i = \{i\}$ for all i , and suppose it is optimal for nodes 3 and 4 to be consecutive on the tour. The chromosomes

1.1 1.9 1.3 1.4 1.8

1.4 1.8 1.5 1.7 1.2

represent the tours $1 \rightarrow 3 \rightarrow 4 \rightarrow 5 \rightarrow 2$ and $5 \rightarrow 1 \rightarrow 3 \rightarrow 4 \rightarrow 2$, respectively. Their offspring might be

1.4 1.9 1.3 1.7 1.8

representing the tour $3 \rightarrow 1 \rightarrow 4 \rightarrow 5 \rightarrow 2$; nodes 3 and 4 are split. If, however, the second solution had been rotated so that it began at node 1, its encoding would be

1.2 1.7 1.4 1.5 1.8

If this individual were to mate with the first one, it would produce offspring in which the genes for

nodes 3 and 4 were no more than 0.2 apart from one another, making it more likely that these nodes would be consecutive.

To avoid rotations, we artificially set the fractional part of the gene for V_1 to 0, ensuring that set 1 will be at the beginning of the tour. To avoid reflections, we must choose between the two orderings of a given tour that begin with cluster 1. The node from V_1 has two neighbors; we require the lower-indexed neighbor to appear in slot 2 on the tour (and the higher-indexed neighbor to appear in slot m). These rules ensure that each tour is stored in a unique manner and that quasi-duplicates are eliminated.

3.6. Termination criteria

The algorithm terminates when 100 iterations have been executed, or when 10 consecutive iterations have failed to improve the best-known solution. These numbers are parameters that may easily be changed. (In our computational tests, the first criterion never held—the GA always terminated because 10 consecutive iterations had failed to improve the solution. We tested termination criteria greater than 10 on our test problems and found substantially longer run times with little improvement in solution quality.)

4. Computational results

4.1. Experimental design

Fischetti et al. [14] describe a branch-and-cut algorithm to solve the symmetric GTSP. In their paper, they derive test problems by applying a partitioning method to 46 standard TSP instances from the TSPLIB library [31]. They provide optimal objective values for each of the problems, and since the partitioning method is deterministic and is described in enough detail to reproduce their test problems, we were able to apply our heuristic to the same problems and compare the results.

We tested our algorithm on all of the problems discussed in [14] except the five problems that use

great-circle (geographical) distances. The problems we tested have between 48 and 442 nodes; m is set to $\lceil n/5 \rceil$ in each case, where n is the number of nodes and m is the number of clusters. Most of the problems use Euclidean distances, but a few have explicit distance matrices given,¹ and one uses a “modified Euclidean” distance measure (see [31] for details); all distance matrices are symmetric. For each problem, we ran the GA five times to examine the algorithm’s performance and its variation from trial to trial. The algorithm was implemented in C++ and tested on a Gateway Profile 4MX desktop computer with a Pentium IV 3.2 GHz processor and 1.0 GB RAM, running under Windows XP.

4.2. Solution quality

Table 1 summarizes the results for each of the problems. The columns are as follows:

Problem. The name of the test problem. The digits at the beginning of the name give the number of clusters (m); those at the end give the number of nodes (n).

Opt Obj Val. The optimal objective value for the problem, as reported in Table 1 of [14].

Opt. The number of trials, out of five, in which the GA found the optimal solution.

Best. The number of trials, out of five, in which the GA found its best solution. (For problems in which an optimal solution was found, this is equal to the “# Opt” column.)

Mean, Minimum, Maximum. The mean, minimum, and maximum objective values returned in the five trials (in the **Value** column), and the respective percentages above the optimal value (in the **Pct** column).

The GA found optimal solutions in at least one of the five trials for 35 of the 41 problems tested (85%). For 32 (78%) of the problems, the

¹ The explicit distance matrices do not satisfy the triangle inequality. Nevertheless, the optimal solutions to the GTSP for these problems contain exactly one node per set, allowing us to use the results from [14].

Table 1
Genetic algorithm results: objective values

Problem	Opt Obj Val	# Opt	# Best	Mean		Minimum		Maximum	
				Value	Pct (%)	Value	Pct (%)	Value	Pct (%)
10ATT48	5394	5	5	5394	0.00	5394	0.00	5394	0.00
10GR48	1834	5	5	1834	0.00	1834	0.00	1834	0.00
10HK48	6386	5	5	6386	0.00	6386	0.00	6386	0.00
11EIL51	174	5	5	174	0.00	174	0.00	174	0.00
12BRAZIL58	15,332	5	5	15,332	0.00	15,332	0.00	15,332	0.00
14ST70	316	5	5	316	0.00	316	0.00	316	0.00
16EIL76	209	5	5	209	0.00	209	0.00	209	0.00
16PR76	64,925	5	5	64,925	0.00	64,925	0.00	64,925	0.00
20KROA100	9711	5	5	9711	0.00	9711	0.00	9711	0.00
20KROB100	10,328	5	5	10,328	0.00	10,328	0.00	10,328	0.00
20KROC100	9554	5	5	9554	0.00	9554	0.00	9554	0.00
20KROD100	9450	5	5	9450	0.00	9450	0.00	9450	0.00
20KROE100	9523	5	5	9523	0.00	9523	0.00	9523	0.00
20RAT99	497	5	5	497	0.00	497	0.00	497	0.00
20RD100	3650	5	5	3650	0.00	3650	0.00	3650	0.00
21EIL101	249	5	5	249	0.00	249	0.00	249	0.00
21LIN105	8213	5	5	8213	0.00	8213	0.00	8213	0.00
22PR107	27,898	5	5	27,898	0.00	27,898	0.00	27,898	0.00
24GR120	2769	5	5	2769	0.00	2769	0.00	2769	0.00
25PR124	36,605	5	5	36,605	0.00	36,605	0.00	36,605	0.00
26BIER127	72,418	5	5	72,418	0.00	72,418	0.00	72,418	0.00
28PR136	42,570	5	5	42,570	0.00	42,570	0.00	42,570	0.00
29PR144	45,886	5	5	45,886	0.00	45,886	0.00	45,886	0.00
30KROA150	11,018	5	5	11,018	0.00	11,018	0.00	11,018	0.00
30KROB150	12,196	5	5	12,196	0.00	12,196	0.00	12,196	0.00
31PR152	51,576	5	5	51,576	0.00	51,576	0.00	51,576	0.00
32U159	22,664	5	5	22,664	0.00	22,664	0.00	22,664	0.00
39RAT195	854	5	5	854	0.00	854	0.00	854	0.00
40D198	10,557	5	5	10,557	0.00	10,557	0.00	10,557	0.00
40KROA200	13,406	5	5	13,406	0.00	13,406	0.00	13,406	0.00
40KROB200	13,111	4	4	13,112	0.00	13,111	0.00	13,114	0.02
45TS225	68,340	4	4	68,352	0.02	68,340	0.00	68,400	0.09
46PR226	64,007	5	5	64,007	0.00	64,007	0.00	64,007	0.00
53GIL262	1013	0	1	1021	0.75	1014	0.10	1025	1.18
53PR264	29,549	5	5	29,549	0.00	29,549	0.00	29,549	0.00
60PR299	22,615	0	1	22,639	0.11	22,620	0.02	22,677	0.27
64LIN318	20,765	2	2	20,894	0.62	20,765	0.00	21,026	1.26
80RD400	6361	0	1	6436	1.19	6416	0.86	6448	1.37
84FL417	9651	0	1	9656	0.05	9654	0.03	9658	0.07
88PR439	60,099	0	1	60,258	0.27	60,100	0.00	60,492	0.65
89PCB442	21,657	0	1	22,026	1.70	21,941	1.31	22,131	2.19

GA found the optimal solution in *every* trial. The GA solved 39 (95%) of the problems to within 1% of optimality on average, and in no case did it return a solution more than 2.2% above optimal. The heuristic tended to return consistent solutions (i.e., the same from trial to trial) for the smaller problems; for larger problems, it returned more

varying solutions (but close to each other in objective value).

4.3. Computation times

Table 2 gives information about running times for each of the trials. The columns are as follows:

Table 2

Genetic algorithm results: running times

Problem	B&C Time	CPU Time			# Iterations			# Impr Iter		
		Mean	Min	Max	Mean	Min	Max	Mean	Min	Max
10ATT48	2.1	0.2	0.0	0.5	12.0	12	12	1.0	1	1
10GR48	1.9	0.3	0.2	0.5	12.6	12	14	1.4	1	2
10HK48	3.8	0.4	0.2	0.9	12.0	12	12	1.0	1	1
11EIL51	2.9	0.2	0.1	0.3	12.0	12	12	1.0	1	1
12BRAZIL58	3.0	0.2	0.1	0.3	13.4	12	16	1.8	1	3
14ST70	7.3	0.2	0.2	0.3	12.0	12	12	1.0	1	1
16EIL76	9.4	0.2	0.2	0.2	12.0	12	12	1.0	1	1
16PR76	12.9	0.2	0.2	0.3	12.8	12	14	1.6	1	2
20KROA100	18.4	0.4	0.3	0.5	12.0	12	12	1.0	1	1
20KROB100	22.2	0.4	0.2	0.5	12.0	12	12	1.0	1	1
20KROC100	14.4	0.3	0.2	0.4	12.0	12	12	1.0	1	1
20KROD100	14.3	0.4	0.2	0.8	12.0	12	12	1.0	1	1
20KROE100	13.0	0.6	0.3	0.8	12.0	12	12	1.0	1	1
20RAT99	51.5	0.5	0.3	0.7	12.0	12	12	1.0	1	1
20RD100	16.6	0.5	0.3	1.0	13.6	12	15	2.2	1	4
21EIL101	25.6	0.4	0.2	0.5	12.2	12	13	1.2	1	2
21LIN105	16.4	0.5	0.3	0.7	12.2	12	13	1.2	1	2
22PR107	7.4	0.4	0.3	0.5	14.4	12	17	1.8	1	2
24GR120	41.9	0.6	0.3	1.0	12.8	12	15	1.4	1	2
25PR124	25.9	0.8	0.6	1.5	13.0	12	15	1.8	1	3
26BIER127	23.6	0.4	0.4	0.5	14.0	12	17	2.0	1	3
28PR136	43.0	0.5	0.3	0.7	13.4	12	14	1.8	1	2
29PR144	8.2	1.0	0.3	2.1	14.6	12	21	2.4	1	4
30KROA150	100.3	0.7	0.3	1.3	13.4	12	15	2.0	1	3
30KROB150	60.6	0.9	0.3	1.2	13.4	12	15	2.0	1	3
31PR152	94.8	1.2	0.9	1.5	15.2	12	17	2.6	1	4
32U159	146.4	0.8	0.4	1.3	12.2	12	13	1.2	1	2
39RAT195	245.9	1.0	0.7	1.4	17.8	13	27	3.8	2	6
40D198	763.1	1.6	1.1	2.7	15.8	13	18	3.0	2	4
40KROA200	187.4	1.8	1.1	2.7	14.4	13	16	2.6	2	4
40KROB200	268.5	1.9	1.4	2.9	22.4	12	35	4.0	1	6
45TS225	37,875.9	2.1	1.4	2.6	24.2	18	30	5.2	4	6
46PR226	106.9	1.5	0.8	2.4	13.0	12	14	1.6	1	2
53GIL262	6624.1	1.9	0.7	3.1	19.4	13	30	3.4	2	6
53PR264	337.0	2.1	1.3	3.5	14.0	13	15	2.4	1	3
60PR299	812.8	3.2	1.6	6.1	26.0	16	42	5.0	3	7
64LIN318	1671.9	3.5	2.4	4.9	20.4	12	33	3.8	1	6
80RD400	7021.4	5.9	3.5	8.9	20.8	17	29	5.2	4	7
84FL417	16,719.4	5.3	2.4	8.6	19.4	15	23	4.4	3	6
88PR439	5422.8	9.5	5.3	12.9	22.8	16	32	5.0	4	6
89PCB442	58,770.5	9.0	4.5	14.5	20.4	15	25	4.0	3	7

Problem. As in Table 1.

CPU Time. The mean, minimum, and maximum CPU time, in seconds, for each of the five trials of the GA.

Iterations. The mean, minimum, and maximum number of iterations before the GA terminated.

Impr Iter. The mean, minimum, and maximum number of iterations during which the GA found a new solution better than the current best solution (excluding the first iteration).

The heuristic executes extremely quickly, with a mean CPU time of less than 10 seconds for all

problems and a maximum of less than 15 seconds. The GA generally found its best solution in the first 2 or 3 iterations or so and spent most of the iterations waiting to terminate. In general, the bulk of the running time (ranging from 65% for the smaller problems to 95% for the larger problems) is spent in the improvement heuristic, suggesting that a more efficient improvement method (such as Renaud and Boctor's G2-opt or G3-opt) may improve the speed of the GA even more.

Note that the GA seems to perform equally well for problems with modified Euclidean distances (10ATT48), and explicit distances (10GR48, 10HK48, 12BRAZIL5, and 24GR120) as it does for those with Euclidean distances. We were unable to test problems with non-geographic clusters (for example, in which nodes are grouped randomly) since no optimal objective values have been published for such problems.

4.4. Comparison to other algorithms

Table 3 compares the performance of the GA with that of several other algorithms (four heuristics and one exact algorithm) on the same TSPLIB problems. The first is the GI^3 heuristic proposed by Renaud and Boctor [32]; the second is Noon's generalized nearest neighbor heuristic [27] combined with the improvement phase from the GI^3 heuristic (results cited in [32]). Renaud and Boctor omit the great-circle problems, as well as several others. The heuristics labeled "FST-Lagr" and "FST-Root" in Table 3 are the Lagrangian procedure and the root-node procedure described in [14]; these procedures produce initial bounds at the root node of the branch-and-cut tree. Table 3 lists both the solution quality (percent above optimal) and CPU time for all five heuristics, as well as the solution time for the branch-and-cut procedure (an exact algorithm, not a heuristic) in [14]. The solution value and times reported for our GA heuristic are for the *first* trial (of the five trials reported in Tables 1 and 2) for each problem, so that we are comparing a single run of our procedure with a single run of each of the others. Minimum values for the "Pct" field in each row are marked in bold. Problems for which the GA did not find the best

solution (among all heuristics) in the first trial but did find it in one of the other trials are indicated with an asterisk (*).

The columns are as follows:

Problem. As in Table 1.

GA. The percentage error of the solution returned in the first of five GA trials and the corresponding CPU time (in seconds) on a Gateway Profile 4MX with Pentium IV 3.2 GHz processor and 1 GB RAM.

GI^3 . The percentage error of the solution returned by the GI^3 heuristic and the corresponding CPU time (in seconds) on a Sun Sparc Station LX, as reported in Table 3 of [32].

NN. The percentage error of the solution returned by the NN heuristic (followed by the improvement phase of GI^3) and the corresponding CPU time (in seconds) on a Sun Sparc Station LX, as reported in Table 3 of [32].

FST-Lagr. The percentage error of the solution returned by Fischetti, Salazar-González, and Toth's Lagrangian heuristic and the corresponding CPU time (in seconds) on an HP 9000/720, as reported in Table 1 of [14].

FST-Root. The percentage error of the solution returned by Fischetti, Salazar-González, and Toth's root-node heuristic and the corresponding CPU time (in seconds) on an HP 9000/720, as reported in Table 1 of [14].

B&C. The CPU time (in seconds) of the branch-and-cut algorithm on an HP 9000/720, as reported in Table 2 of [14].

The first trial of the GA found the best solution among all five heuristics (allowing for ties) in 36 out of the 41 problems tested (88%). The best solution was found in a trial other than the first for two other problems. Two-sided paired *t*-tests confirm that the objective values of the solution returned by the GA are statistically smaller than those returned by GI^3 ($P = 0.002$), NN ($P = 0.001$), and FST-Lagr ($P = 0.005$). The GA also outperforms FST-Root on average, but a two-sided *t*-test could not prove a significant difference ($P = 0.685$). (FST-Root seems to be a much slower procedure than the GA, even allowing for differences in CPU speeds.)

Table 3
Genetic algorithm vs. other algorithms

Problem	GA		GI ³		NN		FST-Lagr		FST-Root		B&C
	Pct	Time	Pct	Time	Pct	Time	Pct	Time	Pct	Time	Time
10ATT48	0.00	0.0					0.00	0.9	0.00	2.1	2.1
10GR48	0.00	0.5					0.00	0.5	0.00	1.9	1.9
10HK48	0.00	0.2					0.00	1.1	0.00	3.8	3.8
11EIL51	0.00	0.1	0.00	0.3	0.00	0.4	0.00	0.4	0.00	2.9	2.9
12BRAZIL58	0.00	0.3					0.00	1.4	0.00	3.0	3.0
14ST70	0.00	0.2	0.00	1.7	0.00	0.8	0.00	1.2	0.00	7.3	7.3
16EIL76	0.00	0.2	0.00	2.2	0.00	1.1	0.00	1.4	0.00	9.4	9.4
16PR76	0.00	0.2	0.00	2.5	0.00	1.9	0.00	0.6	0.00	12.9	12.9
20RAT99	0.00	0.7	0.00	5.0	0.00	7.3	0.00	3.1	0.00	51.4	51.5
20KROA100	0.00	0.4	0.00	6.8	0.00	3.8	0.00	2.4	0.00	18.3	18.4
20KROB100	0.00	0.4	0.00	6.4	0.00	2.4	0.00	3.1	0.00	22.1	22.2
20KROC100	0.00	0.3	0.00	6.5	0.00	6.3	0.00	2.2	0.00	14.3	14.4
20KROD100	0.00	0.4	0.00	8.6	0.00	5.6	0.00	2.5	0.00	14.2	14.3
20KROE100	0.00	0.8	0.00	6.7	0.00	2.8	0.00	0.9	0.00	12.9	13.0
20RD100	0.00	0.3	0.08	7.3	0.08	8.3	0.08	2.6	0.00	16.5	16.6
21EIL101	0.00	0.2	0.40	5.2	0.40	3.0	0.00	1.7	0.00	25.5	25.6
21LIN105	0.00	0.3	0.00	14.4	0.00	3.7	0.00	2.0	0.00	16.2	16.4
22PR107	0.00	0.4	0.00	8.7	0.00	5.2	0.00	2.1	0.00	7.3	7.4
24GR120	0.00	0.5					1.99	4.9	0.00	41.8	41.9
25PR124	0.00	0.6	0.43	12.2	0.00	12.0	0.00	3.7	0.00	25.7	25.9
26BIER127	0.00	0.5	5.55	36.1	9.68	7.8	0.00	11.2	0.00	23.3	23.6
28PR136	0.00	0.5	1.28	12.5	5.54	9.6	0.82	7.2	0.00	42.8	43.0
29PR144	0.00	0.3	0.00	16.3	0.00	11.8	0.00	2.3	0.00	8.0	8.2
30KROA150	0.00	1.3	0.00	17.8	0.00	22.9	0.00	7.6	0.00	100.0	100.3
30KROB150	0.00	1.0	0.00	14.2	0.00	20.1	0.00	9.9	0.00	60.3	60.6
31PR152	0.00	1.5	0.47	17.6	1.80	10.3	0.00	9.6	0.00	51.4	94.8
32U159	0.00	0.6	2.60	18.5	2.79	26.5	0.00	10.9	0.00	139.6	146.4
39RAT195	0.00	0.7	0.00	37.2	1.29	86.0	1.87	8.2	0.00	245.5	245.9
40D198	0.00	1.2	0.60	60.4	0.60	118.8	0.48	12.0	0.00	762.5	763.1
40KROA200	0.00	2.7	0.00	29.7	5.25	53.0	0.00	15.3	0.00	183.3	187.4
40KROB200	0.00	1.4	0.00	35.8	0.00	135.2	0.05	19.1	0.00	268.0	268.5
45TS225	0.00	2.4	0.61	89.0	0.00	117.8	0.09	19.4	0.09	1298.4	37,875.9
46PR226	0.00	1.0	0.00	25.5	2.17	67.6	0.00	14.6	0.00	106.2	106.9
53GIL262	0.79	1.9	5.03	115.4	1.88	122.7	3.75	15.8	0.89	1443.5	6624.1
53PR264	0.00	1.3	0.36	64.4	5.73	147.2	0.33	24.3	0.00	336.0	337.0
60PR299	0.02	6.1	2.23	90.3	2.01	281.8	0.00	33.2	0.00	811.4	812.8
64LIN318	0.00	3.5	4.59	206.8	4.92	317.0	0.36	52.5	0.36	847.8	1671.9
80RD400	1.37*	3.5	1.23	403.5	3.98	1137.1	3.16	59.8	2.97	5031.5	7021.4
84FL417	0.07	2.4	0.48	427.1	1.07	1341.0	0.13	77.2	0.00	16,714.4	16,719.4
88PR439	0.23*	9.1	3.52	611.0	4.02	1238.9	1.42	146.6	0.00	5418.9	5422.8
89PCB442	1.31	10.1	5.91	567.7	0.22	838.4	4.22	78.8	0.29	5353.9	58,770.5

A direct comparison of solution times is difficult since the procedures were coded and tested on different machines. If we make the assumption (conservatively, we feel—see [10]) that the machine used to test the GA is 50 times faster than the Sun Sparc Station LX used to test GI³ and NN and 100 times faster than the HP 9000/720 used

to test FST-Lagr, FST-Root, and the branch-and-cut algorithm, our times are competitive with those of GI³, NN, FST-Root, and branch-and-cut, and in some cases greatly outperform them. The adjusted run times for the GA are longer than those for FST-Lagr, but the solutions produced by the GA are at least as good as those produced

by FST-Lagr for all but one problem. Aside from solution quality and run time, FST-Lagr and the GA each have their advantages: FST-Lagr provides a lower bound on the optimal objective value, while the GA is simpler to code and can easily be modified to incorporate alternate objective functions and constraints.

4.5. Contribution of algorithm features

In this section we examine the relative contribution of each of the features of the GA to the heuristic's overall performance. Table 4 reports the mean cost of the best solution found (out of 5 trials) both before and after the improvement heuristic was

Table 4
Performance of improvement heuristics

Problem	GA value	Pre-Impr cost	% Diff	# 2-Opt	# Swaps	# City Swaps	# Pos Swaps	# City-Pos Swaps
10ATT48	5394.0	9044.4	40.4	3959.4	8010.6	4653.0	826.2	2531.4
10GR48	1834.0	3748.2	51.1	4528.8	8005.2	4387.2	743.0	2875.0
10HK48	6386.0	10,053.6	36.5	4545.4	10,026.6	6343.0	815.6	2868.0
11EIL51	174.0	260.2	33.1	2364.4	3385.2	1694.4	426.0	1264.8
12BRAZIL58	15,332.0	25,461.8	39.8	2615.8	3447.8	1409.8	713.4	1324.6
14ST70	316.0	484.4	34.8	2737.0	3999.6	1869.2	484.6	1645.8
16EIL76	209.0	314.6	33.6	2251.6	2935.0	1092.6	500.6	1341.8
16PR76	64,925.0	106,219.4	38.9	2594.2	4286.8	2190.6	589.6	1506.6
20KROA100	9711.0	29,368.2	66.9	3944.4	6667.2	3589.2	746.8	2331.2
20KROB100	10,328.0	27,595.6	62.6	3997.2	7111.2	4045.8	709.4	2356.0
20KROC100	9554.0	27,713.4	65.5	3188.0	5127.6	2662.0	553.8	1911.8
20KROD100	9450.0	27,031.6	65.0	4431.0	7213.0	3715.6	798.6	2698.8
20KROE100	9523.0	28,711.6	66.8	5291.2	9157.2	4811.2	1057.2	3288.8
20RAT99	497.0	845.0	41.2	5504.8	8418.0	3803.6	955.6	3658.8
20RD100	3650.0	6078.0	39.9	4386.4	7213.2	3620.4	880.4	2712.4
21EIL101	249.0	448.6	44.5	3162.0	4319.6	1581.8	773.6	1964.2
21LIN105	8213.0	22,945.4	64.2	4542.2	7704.6	4003.2	939.6	2761.8
22PR107	27,898.0	46,941.4	40.6	3606.0	4735.2	2314.4	615.0	1805.8
24GR120	2769.0	8184.8	66.2	5707.8	8488.0	3779.4	1176.6	3532.0
25PR124	36,605.0	117,303.2	68.8	8044.8	12,951.6	6978.4	1468.4	4504.8
26BIER127	72,418.0	162,846.2	55.5	3227.4	5149.4	2222.6	1114.4	1812.4
28PR136	42,570.0	99,596.2	57.3	4062.6	6239.8	2769.4	923.2	2547.2
29PR144	45,886.0	127,467.4	64.0	9113.0	14,302.4	7879.4	1640.2	4782.8
30KROA150	11,018.0	32,882.6	66.5	5428.2	8605.4	3979.6	1001.4	3624.4
30KROB150	12,196.0	32,971.0	63.0	7082.6	11,641.4	5802.6	1534.4	4304.4
31PR152	51,576.0	118,768.0	56.6	10,866.0	17,024.6	9128.2	2514.2	5382.2
32U159	22,664.0	75,651.4	70.0	6074.0	9875.2	4992.8	1464.0	3418.4
39RAT195	854.0	1060.8	19.5	6739.8	9787.4	3822.4	1221.2	4743.8
40D198	10,557.0	19,425.2	45.7	11,089.2	17,112.4	7768.2	2858.2	6486.0
40KROA200	13,406.0	47,625.2	71.9	12,220.8	20,380.0	10,162.0	2404.8	7813.2
40KROB200	13,111.6	24,748.2	47.0	12,661.6	20,705.8	10,417.8	2282.6	8005.4
45TS225	68,352.0	154,058.8	55.6	13,614.6	19,805.8	9040.6	2484.6	8280.6
46PR226	64,007.0	84,406.6	24.2	9935.4	12,501.6	4899.0	3756.2	3846.4
53GIL262	1020.6	2234.6	54.3	9653.0	13,945.0	5780.0	1929.0	6236.0
53PR264	29,549.0	66,910.4	55.8	12,157.8	17,236.4	7276.0	2647.2	7313.2
60PR299	22,638.8	47,119.6	52.0	12,483.0	19,081.6	8689.0	2627.2	7765.4
64LIN318	20,893.6	58,746.0	64.4	13,654.4	20,418.6	9082.8	3342.4	7993.4
80RD400	6436.4	14,723.0	56.3	15,313.6	24,164.4	10,858.0	2966.2	10,340.2
84FL417	9655.6	15,636.2	38.2	17,842.2	20,976.6	7542.8	6554.2	6879.6
88PR439	60,258.4	120,506.4	50.0	23,230.4	34484.0	15,683.0	5660.6	13,140.4
89PCB442	22,025.8	52,006.6	57.6	22,462.2	33,104.2	14,059.6	4514.6	14,530.0

performed. It also reports the number of 2-opt and swaps, as well as the breakdown of swaps by type: changing which city from a cluster is included on the tour but not the position of the city, changing the position of a city but not which city is included, and changing both. The columns are as follows:

Problem. As in Table 1.

GA Value. The mean objective value (out of 5 trials) returned by the GA, replicated from Table 1.

Pre-Impr Cost. The mean pre-improvement cost (out of 5 trials) of the best solution found by the GA. (That is, the mean pre-improvement cost

Table 5
GA-Only and Improve-Only results

Problem	GA-Improve	GA Only		Improve Only	
		Value	Pct	Value	Pct
10ATT48	5394.0	6313.4	17.04	5396.0	0.04
10GR48	1834.0	2079.4	13.38	1861.4	1.49
10HK48	6386.0	6921.6	8.39	6386.0	0.00
11EIL51	174.0	227.4	30.69	175.0	0.57
12BRAZIL58	15,332.0	19,124.0	24.73	15,458.6	0.83
14ST70	316.0	450.8	42.66	318.4	0.76
16EIL76	209.0	352.0	68.42	213.4	2.11
16PR76	64,925.0	85,385.2	31.51	65,315.6	0.60
20KROA100	9711.0	20,191.0	107.92	9711.8	0.01
20KROB100	10,328.0	18,537.4	79.49	10,362.0	0.33
20KROC100	9554.0	17,871.6	87.06	9554.0	0.00
20KROD100	9450.0	18,477.0	95.52	9484.2	0.36
20KROE100	9523.0	19,787.6	107.79	9644.8	1.28
20RAT99	497.0	1090.0	119.32	506.0	1.81
20RD100	3650.0	7353.4	101.46	3718.4	1.87
21EIL101	249.0	526.4	111.41	249.6	0.24
21LIN105	8213.0	14,559.4	77.27	8241.8	0.35
22PR107	27,898.0	57,724.6	106.91	27,930.2	0.12
24GR120	2769.0	5628.4	103.26	2866.6	3.52
25PR124	36,605.0	82,713.0	125.96	36,835.4	0.63
26BIER127	72,418.0	154,703.2	113.63	76,666.2	5.87
28PR136	42,570.0	112,674.6	164.68	45,421.4	6.70
29PR144	45,886.0	94,969.2	106.97	46,543.0	1.43
30KROA150	11,018.0	31,199.2	183.17	11,029.0	0.10
30KROB150	12,196.0	34,685.2	184.40	12,549.0	2.89
31PR152	51,576.0	118,813.4	130.37	52,772.8	2.32
32U159	22,664.0	59,099.2	160.76	22,889.2	0.99
39RAT195	854.0	2844.2	233.04	901.0	5.50
40D198	10,557.0	26,453.0	150.57	10,704.0	1.39
40KROA200	13,406.0	46,866.4	249.59	14209.4	5.99
40KROB200	13,111.6	47,303.2	260.77	13,543.0	3.29
45TS225	68,352.0	229,495.2	235.75	70,775.4	3.55
46PR226	64,007.0	263,699.0	311.98	66,250.6	3.51
53GIL262	1020.6	4233.6	314.81	1091.4	6.94
53PR264	29,549.0	145,789.4	393.38	30,603.6	3.57
60PR299	22,638.8	110,977.8	390.21	24,520.0	8.31
64LIN318	20,893.6	94,469.2	352.14	22,272.8	6.60
80RD400	6436.4	34,502.2	436.05	6789.4	5.48
84FL417	9655.6	65,025.6	573.45	10,168.4	5.31
88PR439	60,258.4	364,282.4	504.53	66,364.0	10.13
89PCB442	22,025.8	131,711.8	497.99	23,586.0	7.08

of the 5 post-improvement solutions considered in the “GA Value” column.)

% **Diff.** The percentage difference between “GA Value” and “Pre-Impr Cost.”

2-Opts. The mean number of 2-opts (out of 5 trials) performed during the GA’s execution.

Swaps. The mean number of swaps (out of 5 trials) performed during the GA’s execution.

City Swaps. The mean number of swaps (out of 5 trials) that involved changing a cluster’s included city (but not the position of the cluster on the tour).

Pos Swaps. The mean number of swaps (out of 5 trials) that involved changing the position of a cluster on the tour (but not the cluster’s included city).

City-Pos Swaps. The mean number of swaps (out of 5 trials) that involved changing both a cluster’s included city and its position on the tour.

Clearly, the improvement heuristics play a large role in the overall performance of the GA. The quality of the solution returned by the GA is due in large part to the improvement heuristics, and each heuristic is performed many times during the GA’s execution. Moreover, each type of heuristic (2-opt and all three types of swaps) is used with some frequency. The success of the improvement heuristics led us to test their performance in isolation (without the GA), as well as the GA’s performance without the improvement heuristics. The results are reported in Table 5, whose columns are as follows:

Problem. As in Table 1.

GA-Improve. The mean objective value (out of 5 trials) returned by the GA with the improvement heuristics, replicated from Table 1.

GA Only. The mean objective value (out of 5 trials) returned by the GA with no improvement heuristics applied (in the “Value”) column, and the corresponding percentage above the “GA-Improve” value (in the “Pct” column).

Improve Only. The mean objective value (out of 5 trials) found by generating a population of 100 individuals and performing level-II improvement on the best pN of them and

level-I improvement on the rest (in the “Value”) column, and the corresponding percentage above the “GA-Improve” value (in the “Pct” column).

Table 5 confirms that the improvement heuristics are essential for the success of the GA; without them, the solutions found are anywhere from 8% to 573% worse than those of the full GA. The improvement heuristics by themselves attain values within 10% of the full GA with negligible CPU times (generally less than 0.1 seconds, even for the largest problems). This suggests the use of a standalone “build-and-improve” heuristic in which a population is generated and improved and the best solution is taken with no application of the GA. Such a heuristic could not be relied upon to produce solutions that are as close to optimal as those from the hybrid GA, but it would execute very quickly.

5. Conclusions and future research directions

This paper presents a heuristic to solve the generalized traveling salesman problem. The procedure incorporates a local tour improvement heuristic into a random-key genetic algorithm. The algorithm performed quite well when tested on a set of 41 standard problems with known optimal objective values, finding the optimal solution in the majority of cases. Computation times were small (under 15 seconds), and the GA returned fairly consistent results among trials.

Our GA performs competitively with other heuristics that have been published, both in solution quality and computation time. Moreover, our heuristic has two main advantages over others. First, it is quite simple to implement. Second, it can be extended easily to incorporate alternate objective functions and constraints. For example, one could allow $r_i \geq 1$ nodes to be required for set V_i . Thus, we might require the tour to visit two nodes from V_1 , five nodes from V_2 , and so on. This would require r_i genes for set V_i . Crossover would operate on all r_i genes as a group, rather than individually, to maintain feasibility. Or, by including demand-weighted “medial” distances in the

objective function, one could solve the *Median Tour Problem* [8,34]. Similarly, by including secondary tours that connect the customers in each cluster to the main tour one can solve the *Traveling Circus Problem* [34]; this problem would require further modifications to the encoding scheme, but it is particularly appealing because it has as a special case the vehicle routing problem (VRP) and the multi-depot VRP. Of course, the success of our algorithm for the classical GTSP does not imply its success for these variations; further computational tests would need to be performed to judge its applicability to these problems.

References

- [1] J.T. Alander (Ed.), Proceedings of the First Nordic Workshop on Genetic Algorithms and their Applications (INWGA), January 9–12, Vaasa Yliopiston Julkaisuja, Vaasa, 1995.
- [2] W. Banzhaf, J. Daida, A.E. Eiben, M.H. Garzon, V. Honavar, M. Jakiela, R.E. Smith (Eds.), Proceedings of the Genetic and Evolutionary Computation Conference, Orlando, FL, July 13–17, Morgan Kaufmann, San Mateo, CA, 1999.
- [3] J.C. Bean, Genetic algorithms and random keys for sequencing and optimization ORSA, Journal on Computing 6 (1994) 154–160.
- [4] R.K. Belew, L.B. Booker (Eds.), Proceedings of the Fourth International Conference on Genetic Algorithms, University of California, San Diego, July 13–16, Morgan Kaufmann, San Mateo, CA, 1991.
- [5] D. Ben-Arieh, G. Gutin, M. Penn, A. Yeo, A. Zverovitch, Process planning for rotational parts using the generalized traveling salesman problem, International Journal of Production Research 41 (11) (2003) 2581–2596.
- [6] D. Ben-Arieh, G. Gutin, M. Penn, A. Yeo, A. Zverovitch, Transformations of generalized ATSP into ATSP, Operations Research Letters 31 (2003) 357–365.
- [7] A.G. Chentsov, L.N. Korotayeva, The dynamic programming method in the generalized traveling salesman problem, Mathematical and Computer Modelling 25 (1) (1997) 93–105.
- [8] J.R. Current, D.A. Schilling, The median tour and maximal covering tour problems: Formulations and heuristics, European Journal of Operational Research 73 (1994) 114–126.
- [9] V. Dimitrijevic, Z. Saric, An efficient transformation of the generalized traveling salesman problem into the traveling salesman problem on digraphs, Information Sciences 102 (1–4) (1997) 105–110.
- [10] J.J. Dongarra, Performance of various computers using standard linear equations software, Technical Report CS-89-85, Computer Science Department, University of Tennessee, 2004.
- [11] M. Dror, M. Haouari, Generalized Steiner problems and other variants, Journal of Combinatorial Optimization 4 (4) (2000) 415–436.
- [12] M. Dror, A. Langevin, A generalized traveling salesman problem approach to the directed clustered rural postman problem, Transportation Science 31 (2) (1997) 187–192.
- [13] M. Fischetti, J.J. Salazar-González, P. Toth, The symmetrical generalized traveling salesman polytope, Networks 26 (2) (1995) 113–123.
- [14] M. Fischetti, J.J. Salazar-González, P. Toth, A branch-and-cut algorithm for the symmetric generalized traveling salesman problem, Operations Research 45 (3) (1997) 378–394.
- [15] N. Garg, G. Konjevod, R. Ravi, A polylogarithmic approximation algorithm for the group Steiner tree problem, Journal of Algorithms 37 (1) (2000) 66–84.
- [16] A. Henry-Labordere, The record balancing problem—A dynamic programming solution of a generalized traveling salesman problem, Revue Francaise D Informatique De Recherche Operationnelle 3 (NB2) (1969) 43–49.
- [17] J.R. Koza, D.E. Goldberg, D.B. Fogel, R.L. Riolo (Eds.), Genetic Programming: Proceedings of the First Annual Conference, Stanford University, July 28–31, MIT Press, Cambridge, 1996.
- [18] V.M. Kureichick, V.V. Miagkikh, A.P. Topchy, Genetic algorithm for solution of the traveling salesman problem with new features against premature convergence, Working Paper, Taganrog State University of Radio-Engineering, Taganrog, Russia, 1996.
- [19] G. Laporte, A. Asef-Vaziri, C. Sriskandarajah, Some applications of the generalized travelling salesman problem, Journal of the Operational Research Society 47 (12) (1996) 1461–1467.
- [20] G. Laporte, H. Mercure, Y. Nobert, Finding the shortest Hamiltonian circuit through n clusters: A Lagrangian approach, Congressus Numerantium 48 (1985) 277–290.
- [21] G. Laporte, H. Mercure, Y. Nobert, Generalized traveling salesman problem through n -sets of nodes—The asymmetrical case, Discrete Applied Mathematics 18 (2) (1987) 185–197.
- [22] G. Laporte, Y. Nobert, Generalized traveling salesman problem through n -sets of nodes—An integer programming approach, INFOR 21 (1) (1983) 61–75.
- [23] G. Laporte, F. Semet, Computational evaluation of a transformation procedure for the symmetric generalized traveling salesman problem, INFOR 37 (2) (1999) 114–120.
- [24] D. Levine, Application of a hybrid genetic algorithm to airline crew scheduling, Computers & Operations Research 23 (6) (1996) 547–558.
- [25] Y.N. Lien, E. Ma, B.W.S. Wah, Transformation of the generalized traveling-salesman problem into the standard traveling-salesman problem, Information Sciences 74 (1–2) (1993) 177–189.
- [26] K.F. Man, K.S. Tang, S. Kwong, Genetic Algorithms: Concepts and Designs, Springer, New York, 1999.

- [27] C.E. Noon, The generalized traveling salesman problem, PhD thesis, University of Michigan, 1988.
- [28] C.E. Noon, J.C. Bean, A Lagrangian based approach for the asymmetric generalized traveling salesman problem, *Operations Research* 39 (4) (1991) 623–632.
- [29] C.E. Noon, J.C. Bean, An efficient transformation of the generalized traveling salesman problem, *INFOR* 31 (1) (1993) 39–44.
- [30] J.E. Rawlins, Gregory (Eds.), *Foundations of Genetic Algorithms*, Morgan Kaufmann, San Mateo, CA, 1991.
- [31] G. Reinelt, TSPLIB—A traveling salesman problem library, *ORSA Journal on Computing* 4 (1996) 134–143.
- [32] J. Renaud, F.F. Boctor, An efficient composite heuristic for the symmetric generalized traveling salesman problem, *European Journal of Operational Research* 108 (3) (1998) 571–584.
- [33] J. Renaud, F.F. Boctor, G. Laporte, A fast composite heuristic for the symmetric traveling salesman problem, *INFORMS Journal on Computing* 4 (1996) 134–143.
- [34] C.S. Revelle, G. Laporte, The plant location problem: New models and research prospects, *Operations Research* 44 (6) (1996) 864–874.
- [35] J.P. Saskaena, Mathematical model of scheduling clients through welfare agencies, *Journal of the Canadian Operational Research Society* 8 (1970) 185–200.
- [36] P. Slavík, On the approximation of the generalized traveling salesman problem, Working Paper, Department of Computer Science, SUNY-Buffalo, 1997.
- [37] W.M. Spears, K.A. DeJong, On the virtues of parameterized uniform crossover, in: *Proceedings of the Fourth International Conference on Genetic Algorithms*, 1991, pp. 230–236.
- [38] S.S. Srivastava, S. Kumar, R.C. Garg, P. Sen, Generalized traveling salesman problem through n sets of nodes, *Journal of the Canadian Operational Research Society* 7 (1970) 97–101.
- [39] L.D. Whitley (Ed.), *Foundations of Genetic Algorithms* 2, Morgan Kaufmann, San Mateo, CA, 1993.
- [40] G. Winter, J. Périaux, M. Galán, P. Cuesta (Eds.), *Genetic Algorithms in Engineering and Computer Science*, Wiley, New York, 1995.
- [41] A.M.S. Zalzalá, P.J. Fleming (Eds.), *Genetic Algorithms in Engineering Systems*, The Institution of Electrical Engineers, London, 1997.